

System Architecture Overview

■ ■ ■■ tokenize: [SELECT, *, FROM, ...] ■ ■

[illegible]

■ ■ ■ ■

■ ■ Size: Configurable per page ■ ■

[REDACTED]
[REDACTED]

[illegible]

▼

[illegible]

[REDACTED]

■ ■ ■ ■

■ ■ products.json ← Stores products (custom schema) ■ ■

```
■ ■ stores.json ← Stores stores (custom schema) ■ ■
```

■ ■ ... ← One file per table ■ ■

■ ■ ■ ■

■ ■ Read Flow: disk file → bincode deserialization → RAM ■ ■

[REDACTED]

[REDACTED]

[illegible]

...

Flow 1: SELECT Query Execution

■ 5. SORT ■

■ ORDER BY name ■

■ (alphabetical sort) ■

100

■ Results: $[A \rightarrow Z]$ ■

▼

■ 6. LIMIT ■

■ LIMIT 10 ■

■ (take first 10) ■

□ □

■ Results: 2 rows ✓ ■

▼

[illegible]

■ 7. RETURN RESULTS ■

□ □

■ (1, Widget, 19.99, ...) ■

■ (2, Gadget, 29.99, ...) ■

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52

...

Flow 2: INSERT Query Execution

...

User Input:

```
"INSERT INTO products VALUES (3, 'Doohickey', '34.99', '75', 'Misc')"
```


▼

[illegible]

■ 1. TOKENIZE ■ → ["INSERT", "INTO", "products", ...]

▼

■ 2. PARSE ■

■ VALUES clause: ■

[illegible]

▼

■ 3. GET TABLE SCHEMA ■

■ stock, category] ■

[illegible]

▼

[illegible]

■ Maps values→columns ■

 $\square \} \square$

XX

■ ■ ■ Statement::AlterTable { ... }

■ ■ State: Stateless (pure functions)

TABLE.RS (Data Structure & Operations)

■ ■ Row struct:

■ ■ ■ id: u32 (fixed)

■ ■ ■ username: String (fixed)

■ ■ ■ email: String (fixed)

■ ■ ■ extras: HashMap (dynamic)

■ ■ Table struct:

■ ■ ■ schema: Vec (column definitions)

■ ■ ■ rows: Vec (data storage)

■ ■ ■ id_index: BTreeMap (fast id lookups)

■ ■ ■ username_index: BTreeMap (fast username lookups)

■ ■ ■ email_index: BTreeMap (fast email lookups)

■ ■ Operations:

■ ■ ■ insert(&mut; self, row: Row)

■ ■ ■ select_where_complex(&self;,, conditions) → Vec

■ ■ ■ update(&mut; self, id: u32, updates: HashMap)

■ ■ ■ delete(&mut; self, id: u32)

■ ■ ■ get_value(&row;,, column) → Option

■ ■ ■ compute_aggregate(...) → Result

■ ■ Uses: BTreeMap for indexing

PAGER.RS (Storage Management)

■ ■ Page structure:

■ ■ ■ data: Vec (rows in page)

■ ■ ■ id: u64 (unique page id)

■ ■ Pager structure:

■ ■ ■ pages: HashMap

■ ■ ■ current_page_id: u64

■ ■ ■ table_name: String

■ ■ Operations:

■ ■ ■ add_row_to_page(row)

■ ■ ■ get_page(id) → Page

■ ■ ■ flush_to_disk()


```

    extras: {}
  ],
  id_index: BTreeMap {1→row, 2→row},
  username_index: BTreeMap {
    "alice"→[1], "bob"→[2]
  },
  email_index: BTreeMap { ... }
}

"products" → Table {
  schema: [id, name, price, stock,
  category],
  rows: Vec [
    {id:1, user:null, email:null,
    extras: {
      "name": "Widget",
      "price": "19.99",
      "stock": "100",
      "category": "Tools"
    }},
    {id:2, user:null, email:null,
    extras: { ... }}
  ],
  id_index: BTreeMap { 1→row, 2→row },
  username_index: BTreeMap {} (empty),
  email_index: BTreeMap {} (empty)
}

...

---
```

Query Processing Pipeline Detail

...

Input: `SELECT * FROM products WHERE stock > '50' ORDER BY price DESC LIMIT 5`

Stage 1: LEXICAL ANALYSIS

■ ■ Input: Raw SQL string

■ ■ Process: Tokenize into keywords and values

■ ■ Output: `[SELECT, *, FROM, products, WHERE, stock, >, 50, ORDER, BY, ...]`

Stage 2: SYNTAX ANALYSIS

■ ■ Input: Token stream

■ ■ Process: Parse into Statement structure

■ ■ Output: `Statement::Select {`

`table: "products",`

`columns: ["*"],`

`where_clause: Some(Condition),`

`order_by: Some(OrderBy { column: "price", asc: false } ),`

`limit: Some(5)`

`}`

Stage 3: SEMANTIC ANALYSIS

■ ■ Input: Parsed statement

■ ■ Process:

■ ■ ■ Get table from registry

■ ■ ■ Validate columns exist

■ ■ ■ Check data types

■ ■ Output: Validated statement + table metadata

Stage 4: OPTIMIZATION

■ ■ Input: Validated statement + table metadata

■ ■ Process:

■ ■ ■ Check if column is indexed

■ ■ ■ "stock" not indexed → full scan needed

■ ■ ■ Estimate rows to process

■ ■ ■ Choose execution strategy

■ ■ Output: Execution plan (full scan)

Stage 5: EXECUTION

■■ Process: `table.select_where_complex()`

■■ Sub-steps:

■ ■■ Iterate through all rows in table

■ ■■ Evaluate: `row.stock > '50'` for each row

■ ■■ Collect matching rows

■ ■■ Output: Vec with 2 matches

■■ Output: [Widget (stock: 100), Gadget (stock: 50)]

Stage 6: FILTERING/TRANSFORMATION

■■ Process: Apply any transformations

■■ Our query: `SELECT *` (no transformation)

■■ Output: Unchanged rows

Stage 7: SORTING

■■ Process: `ORDER BY price DESC`

■■ Algorithm: In-memory quicksort on price column

■■ Output: [Widget (\$19.99), Gadget (\$29.99)]

Stage 8: LIMITING

■■ Process: `LIMIT 5`

■■ Logic: Take first 5 results (we have 2)

■■ Output: [Widget (\$19.99), Gadget (\$29.99)]

Stage 9: FORMATTING

■■ Process: Convert to display format

■■ Get schema: ["id", "name", "price", "stock", "category"]

■■ For each row, build output string

■■ Output (to console):

(1, Widget, 19.99, 100, Tools)

(2, Gadget, 29.99, 50, Electronics)

...

Index Strategy Visualization

Table: products

Columns: [id, name, price, stock, category]

Data:

Row 1: {id: 1, name: Widget, price: 19.99, stock: 100, category: Tools}

Row 2: {id: 2, name: Gadget, price: 29.99, stock: 50, category: Electronics}

Row 3: {id: 3, name: Doohickey, price: 34.99, stock: 75, category: Misc}

Indexes Created:

- ✓ id_index (B-Tree) ← Fast lookup on id

- ✓ username_index (B-Tree) ← Unused for products table

- ✓ email_index (B-Tree) ← Unused for products table

$\times$ name_index $\leftarrow$ NOT created (optimization)

~~X~~ price_index ← NOT created (optimization)

~~X~~ stock_index $\leftarrow$ NOT created (optimization)

x category_index ← NOT created (optimization)

Index Contents:

id_index:

[illegible]

■ BTreeMap ■

[illegible]

■ $1 \rightarrow (0, 0)$ [points to Row 1] ■

■ $2 \rightarrow (0, 1)$ [points to Row 2] ■

■ $3 \rightarrow (0, 2)$ [points to Row 3] ■

Performance Analysis:

[illegible]

■ Query: SELECT * WHERE id = 2 ■

[illegible]

■ Execution: Look in id index ■

■ Complexity: $O(\log 3) \approx 2$ comparisons ■

■ Result: Found at (page 0, row 1) ■

■ }

Step 3: Table Created

■ ■ New Table {

```
■ schema: ["id", "name", "department", "salary", "hire_date"],
```

■ rows: [],

- id_index: empty,

- username_index: empty,

■ email_index: empty

■ }

Step 4: File Created

■■ employees.json (created with empty array)

Persistent State After Command:

■■ schemas.json → Updated with new schema

■■ employees.json → Empty table file

- In-memory registry → "employees" accessible for next query

...

Error Handling Flow

...

User Input: `SELECT * FROM nonexistent_table WHERE id = 1`

▼

XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX

■ Execute Statement ■

XXXXXXXXXXXXXXXXXXXXXXXXXXXX

■

▼

XX

■ Load Table by Name: ■

```
■ schemas.get("nonexistent") ■
```

■ ■ ■ Heap: ['a', 'l', 'i', 'c', 'e'] ■ ■ ■

- Execution: ■
- 1. Load orders table → [Row1, Row2, Row3] ■
- 2. Get schema: [id, customer_id, product_id, qty, date]■
- 3. Iterate rows: ■
- ■■ Check row.get_value("customer_id") ■
- Row1: "1" == "1" ✓ Include ■
- Row2: "1" == "1" ✓ Include ■
- Row3: "2" != "1" ✗ Skip ■
- 4. Return: [Row1, Row2] ■
- 5. Format output with schema: ■
- (1, 1, 101, 2, 2024-01-15) ■

...

...

[REDACTED]

[REDACTED]

[REDACTED]

[REDACTED]

Summary: Complete System Overview

Three-Tier Architecture:

1. **User Interface Tier**: REPL with SQL command parsing
2. **Business Logic Tier**: Query execution and table operations
3. **Data Tier**: B-Tree indexes and persistent storage

Key Innovation: Dynamic Schemas

- **Problem**: Fixed 3-column schema limited database completeness
- **Solution**: HashMap extras + schema registry
- **Result**: Support for unlimited custom tables with any columns

Performance Characteristics:

- Indexed lookups: $O(\log n)$ - **Excellent**
- Full table scans: $O(n)$ - **Good for typical datasets**
- Sorting: $O(n \log n)$ - **Standard**
- Grouping: $O(n)$ - **Efficient**

Data Flow Summary:

User Input → Parser → Executor → Table Ops → Indexes → Storage → User Output

All layers communicate seamlessly through schema-aware operations.

Last Updated: Current Session

Diagram Version: 1.0

Architecture Status: Stable and Production-Ready